RAG improves observed correctness; SFT gains are mixed

Mean semantic rubric score · 0–100% · labels show scored questions / 76 · whiskers: conditional 95% bootstrap intervals

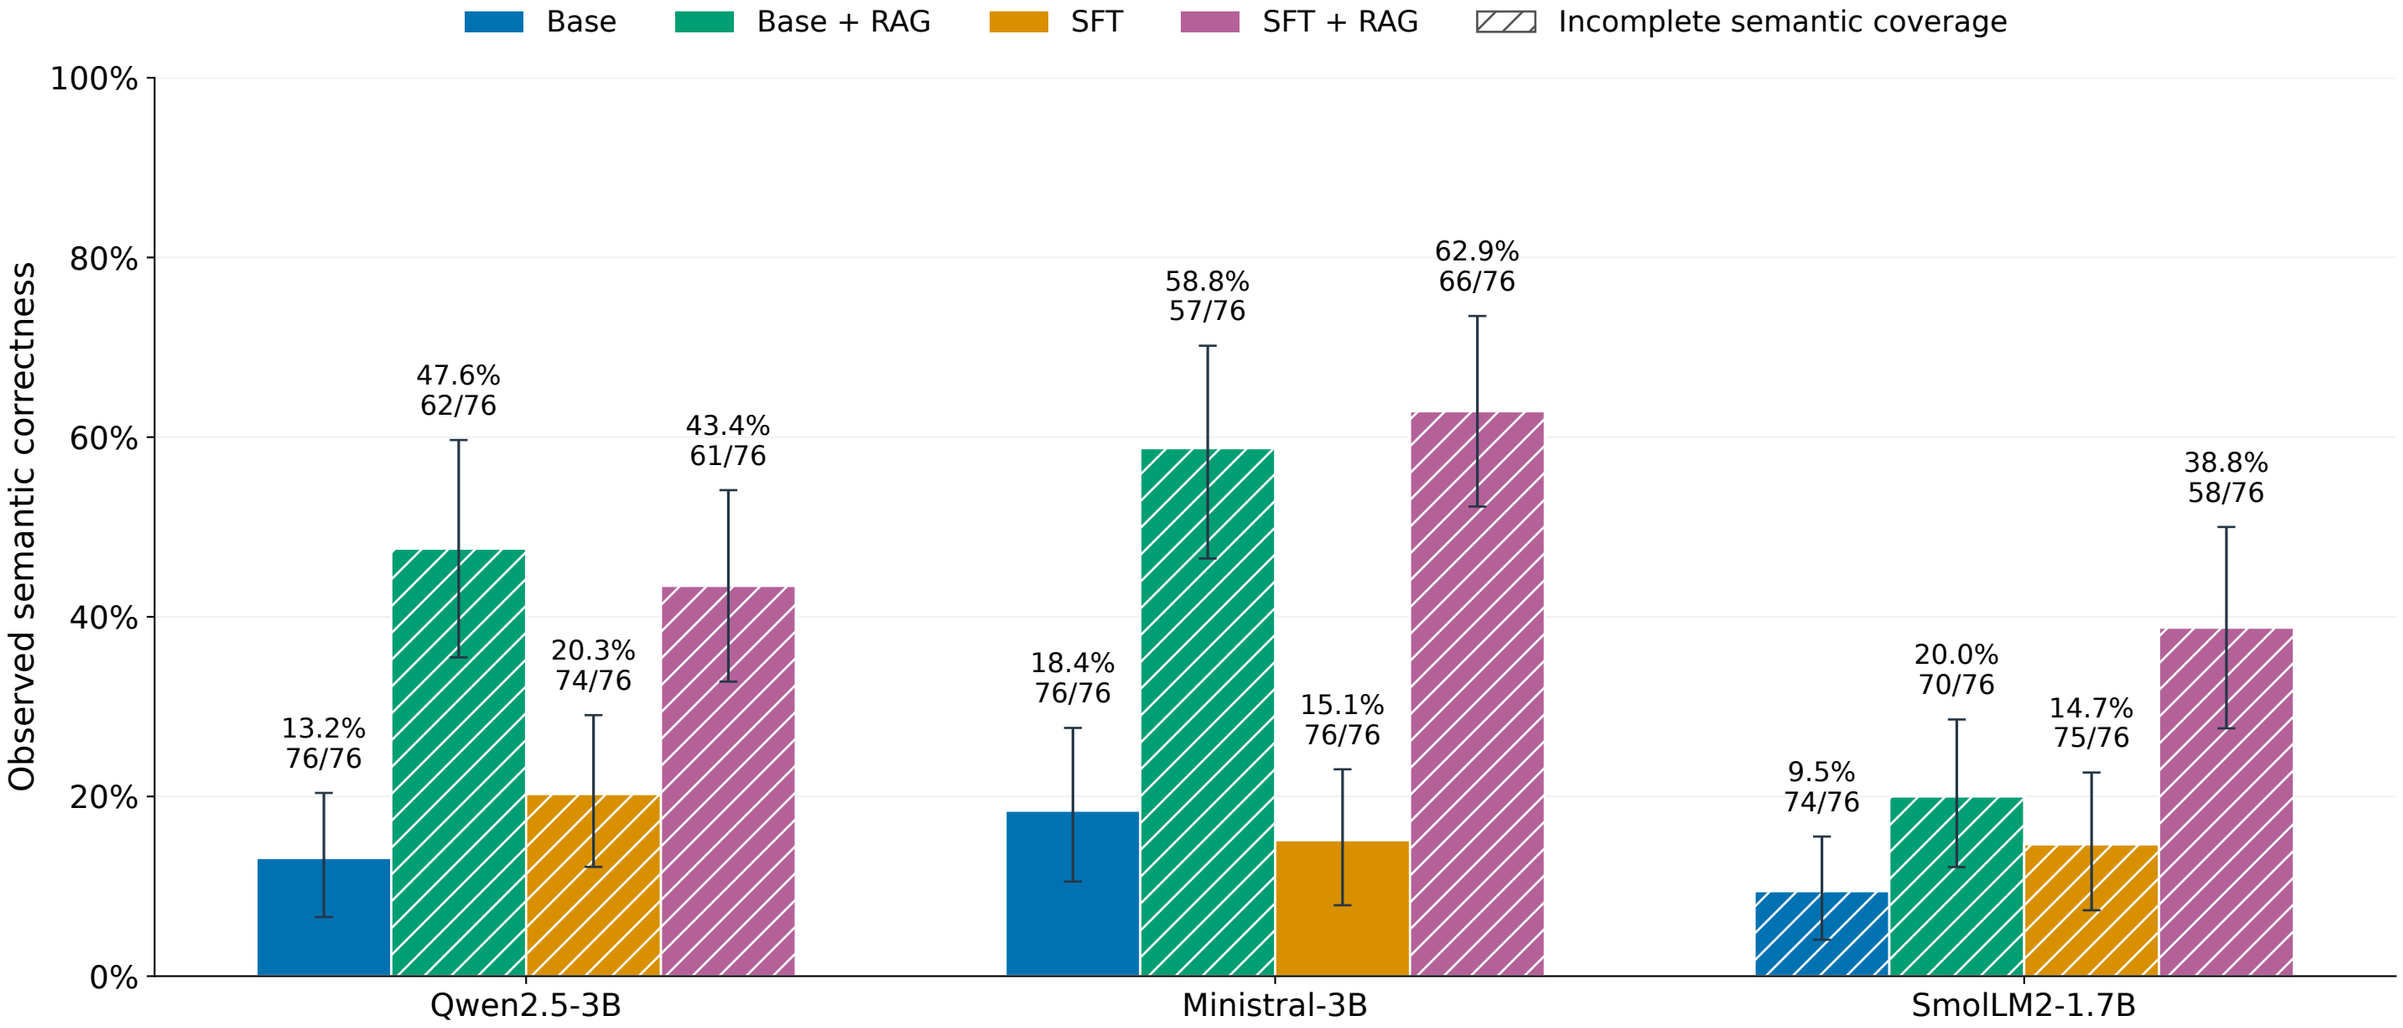

Intervals condition on observed judged questions; they do not include judge error, missingness bias or model-training uncertainty.

PROVISIONAL · private checkpoint v427 · 825/912 judged · mixed-provider single-judge · DeepSeek N/A